

238P Operating Systems, Fall 2018

xv6 Boot Recap: Transitioning from 16 bit mode to 32 bit mode

3 November 2018

Aftab Hussain

University of California, Irvine

BIOS → xv6 Boot loader

what it does

Sets up the hardware.

Transfers control to the Boot Loader.

BIOS → xv6 Boot loader

what it does

Sets up the hardware.

Transfers control to the Boot Loader.

how it transfers control to the Boot Loader

Boot loader is **loaded from** the 1st 512-byte sector of the boot disk.

This 512-byte sector is known as the **boot sector**.

Boot loader is **loaded at** 0x7c00.

Sets processor's **ip** register to 0x7c00.

BIOS → xv6 Boot loader

2 source source files

bootasm.S - 16 and 32 bit assembly code.

bootmain.c - C code.

BIOS → xv6 Boot loader

2 source source files

bootasm.S - 16 and 32 bit assembly code.

bootmain.c - C code.

executing bootasm.S

1. **Disable interrupts** using [cli](#) instruction. ([Code](#)).
 - > Done in case BIOS has setup any of its interrupt handlers were initialized while setting up the hardware. Also, BIOS is not running anymore, so better to disable them.
 - > Clear segment registers. Use xor for %ax, and copy it to the rest ([Code](#)).
2. **Switch from real mode to protected mode.** (References: [a](#), [b](#)).
 - > Note the difference between processor modes and kernel privilege modes
 - > We do the above switch to increase the size of the memory we can address.

BIOS → xv6 Boot loader

executing bootasm.S

Let's expand on this a little bit

2. Switch from real mode to protected mode.

Addressing in Real Mode

In real mode, the processor sends 20-bit addresses to the memory.
(e.g. Intel processors 8086, 8088).

However, it has eight general 16-bit registers + 16 bit segment registers.

How to generate a 20 bit address from a 16 bit register?

Say processor fetches a data read/write instruction.
The processor would then use the data segment register (%ds).

To pass the 16-bit address obtained from %ds
to the memory (which accepts 20-bit addresses) we left
shift the 16-bit register by 4 bits.

This is equivalent to **adding the register
to itself 16 times.**
(try it out in a bit calculator - e.g. gnu calculator)

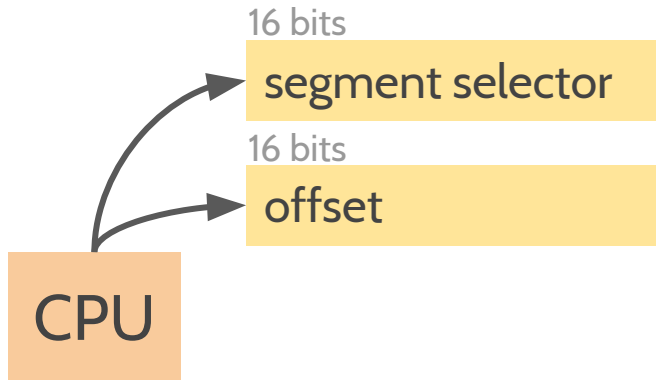
*Use other kinds of
segment registers
for other kinds of instructions.
For example for stack reads/writes use %ss.*

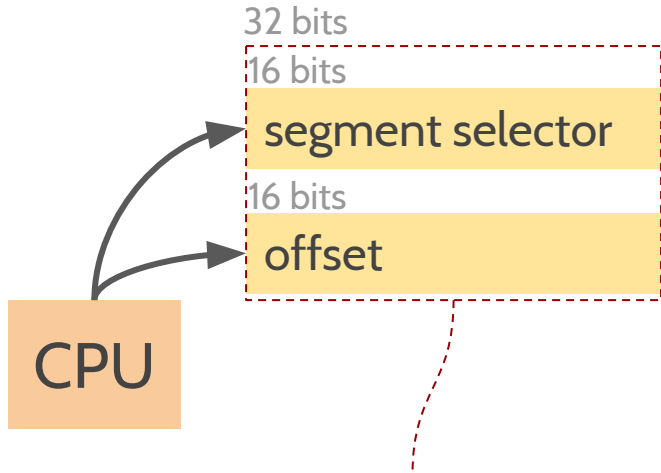
Say processor fetches a data read/write instruction.
The processor would then use the data segment register (%ds).

To pass the 16-bit address obtained from %ds
to the memory (which accepts 20-bit addresses) we left
shift the 16-bit register by 4 bits.

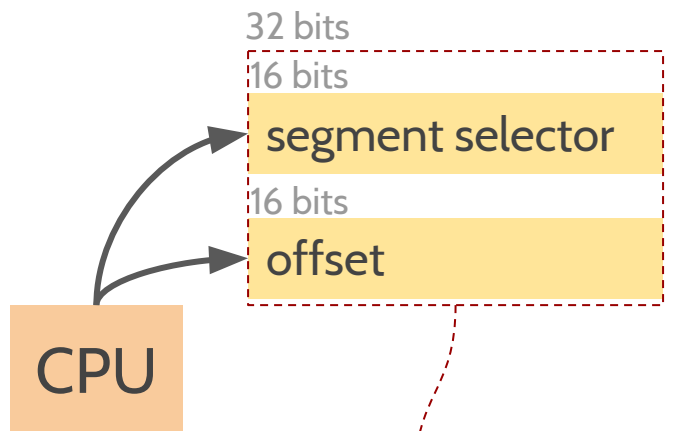
This is equivalent to **adding the register
to itself 16 times.**
(try it out in a bit calculator - e.g. gnu calculator)

So let's look at how the
address translation process takes place in real mode...



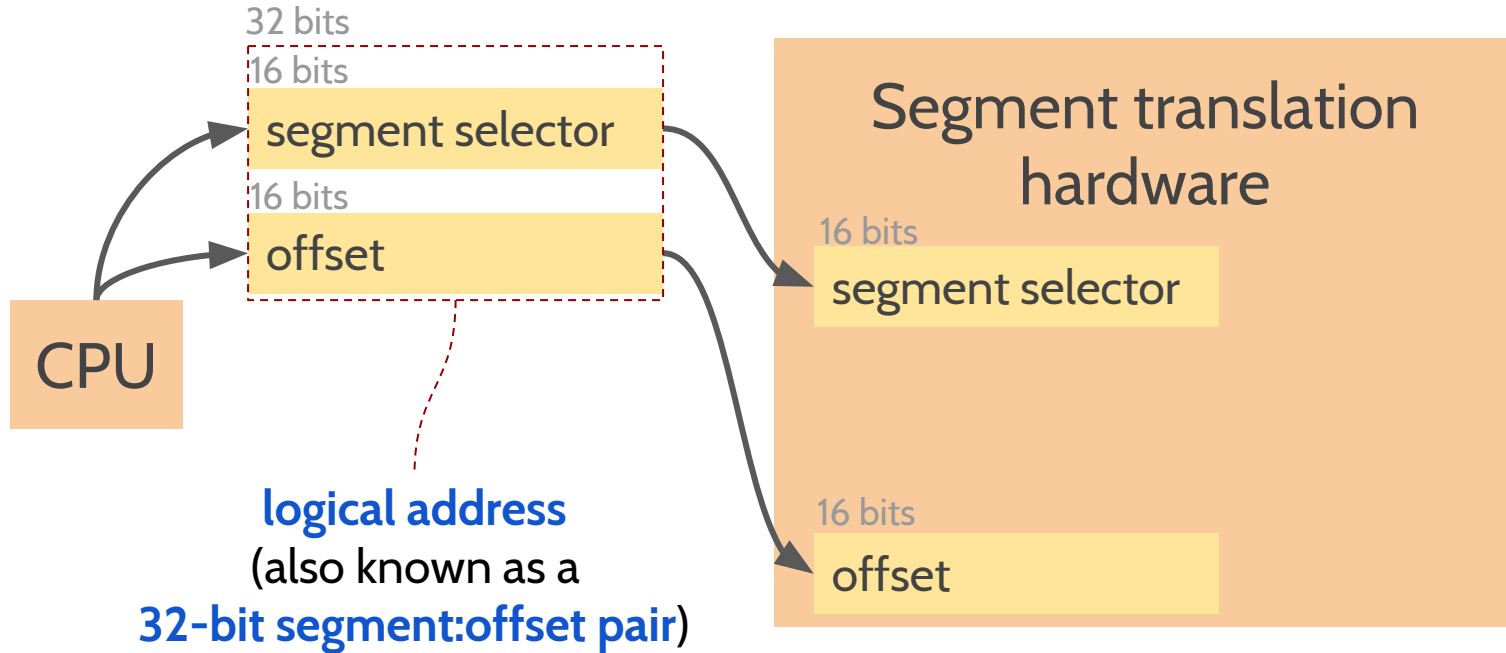


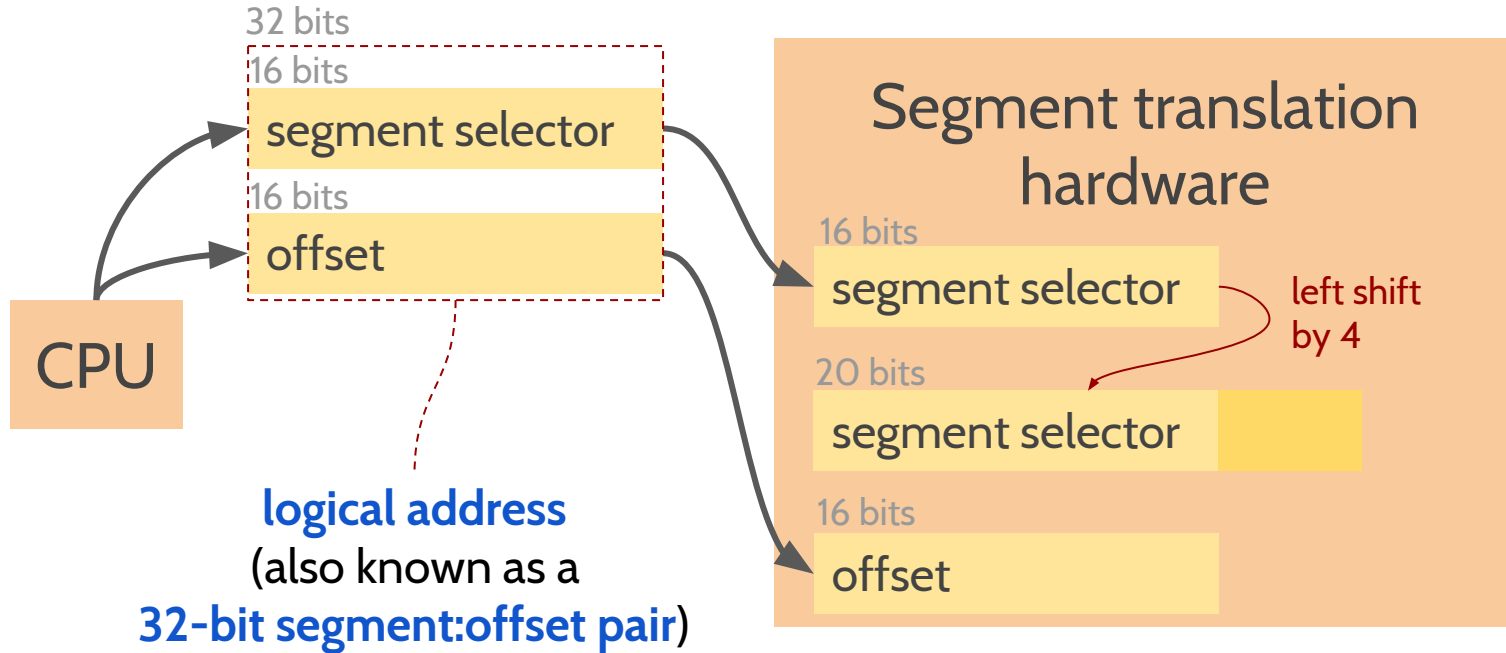
logical address
(also known as a
32-bit segment:offset pair)

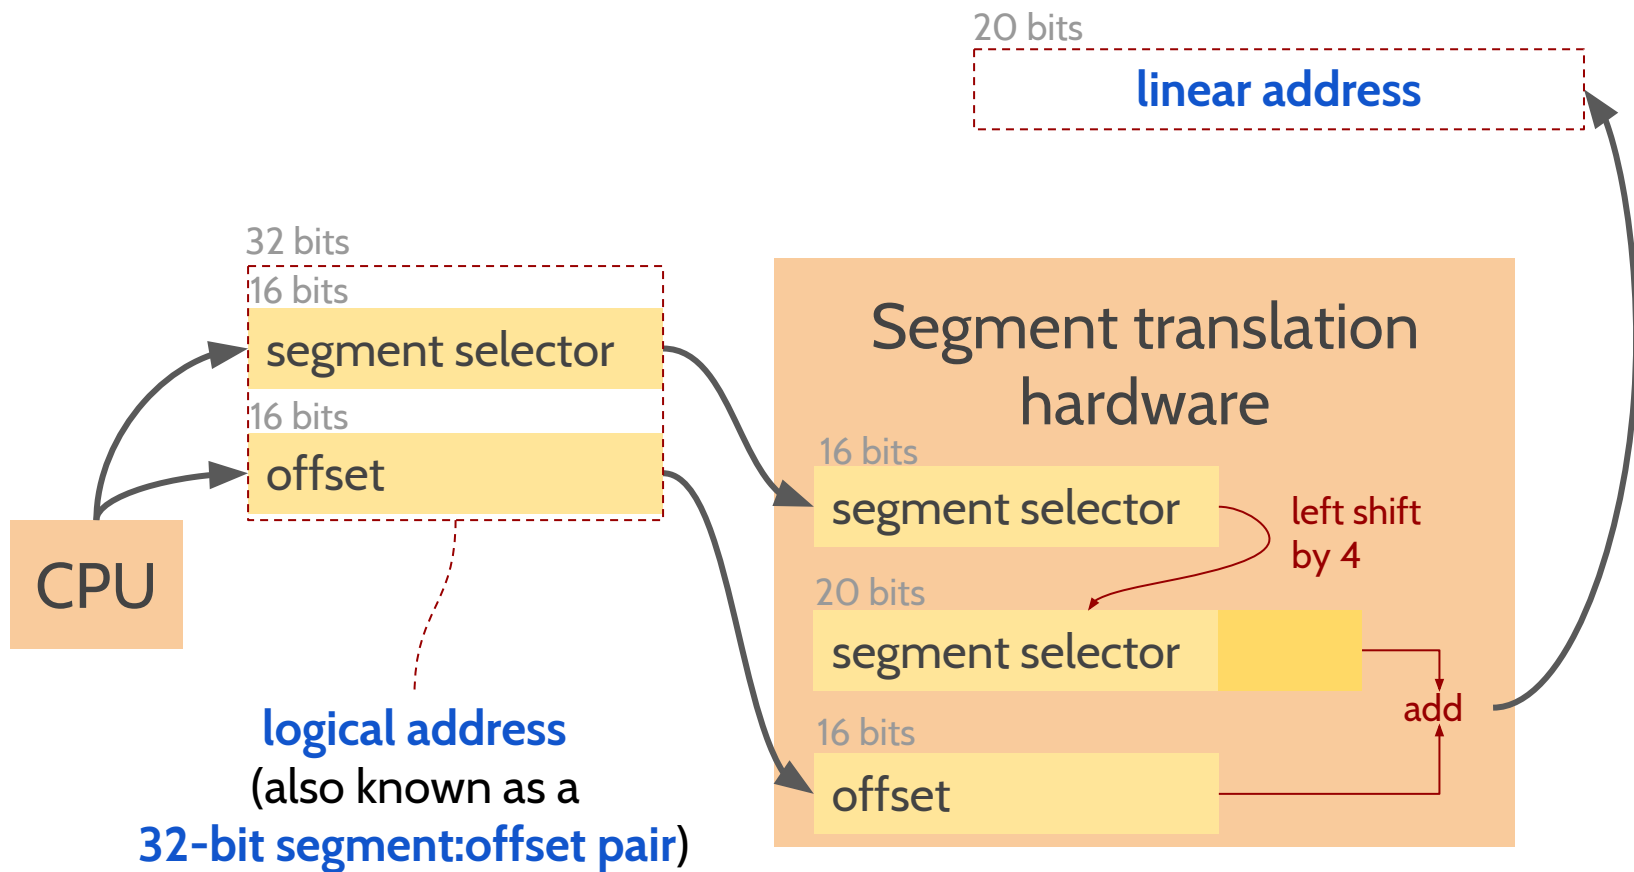


logical address
(also known as a
32-bit segment:offset pair)

xv6 refers to this x86 logical address as a **virtual address**







directly
corresponds to the
physical address

20 bits

linear address

32 bits

16 bits

segment selector

16 bits

offset

CPU

logical address
(also known as a
32-bit segment:offset pair)

Segment translation
hardware

16 bits

segment selector

left shift
by 4

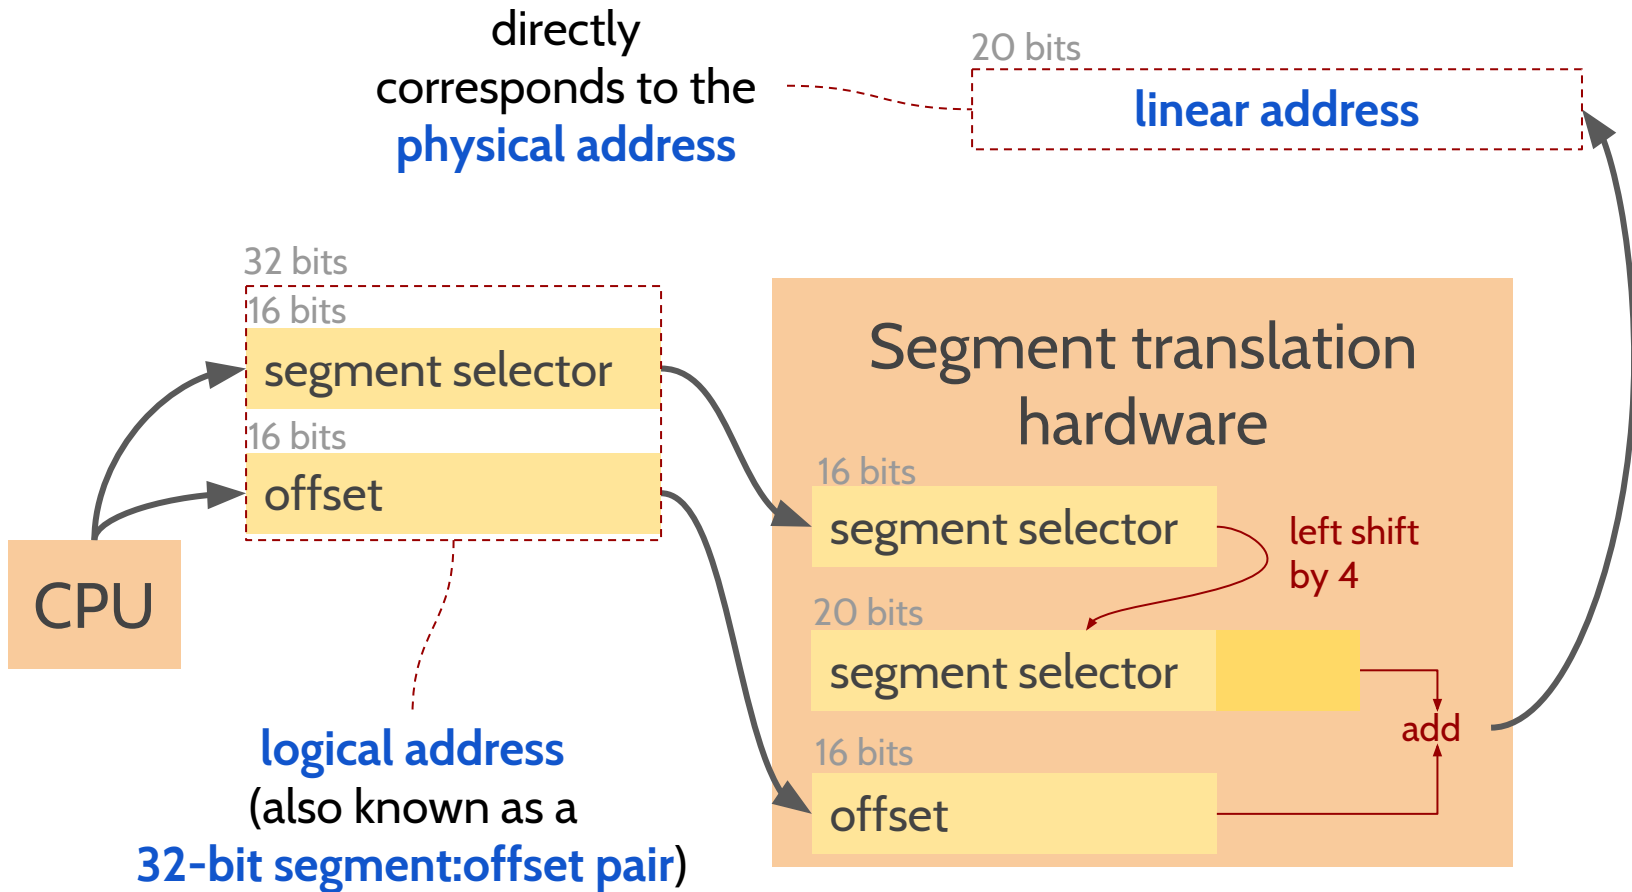
20 bits

segment selector

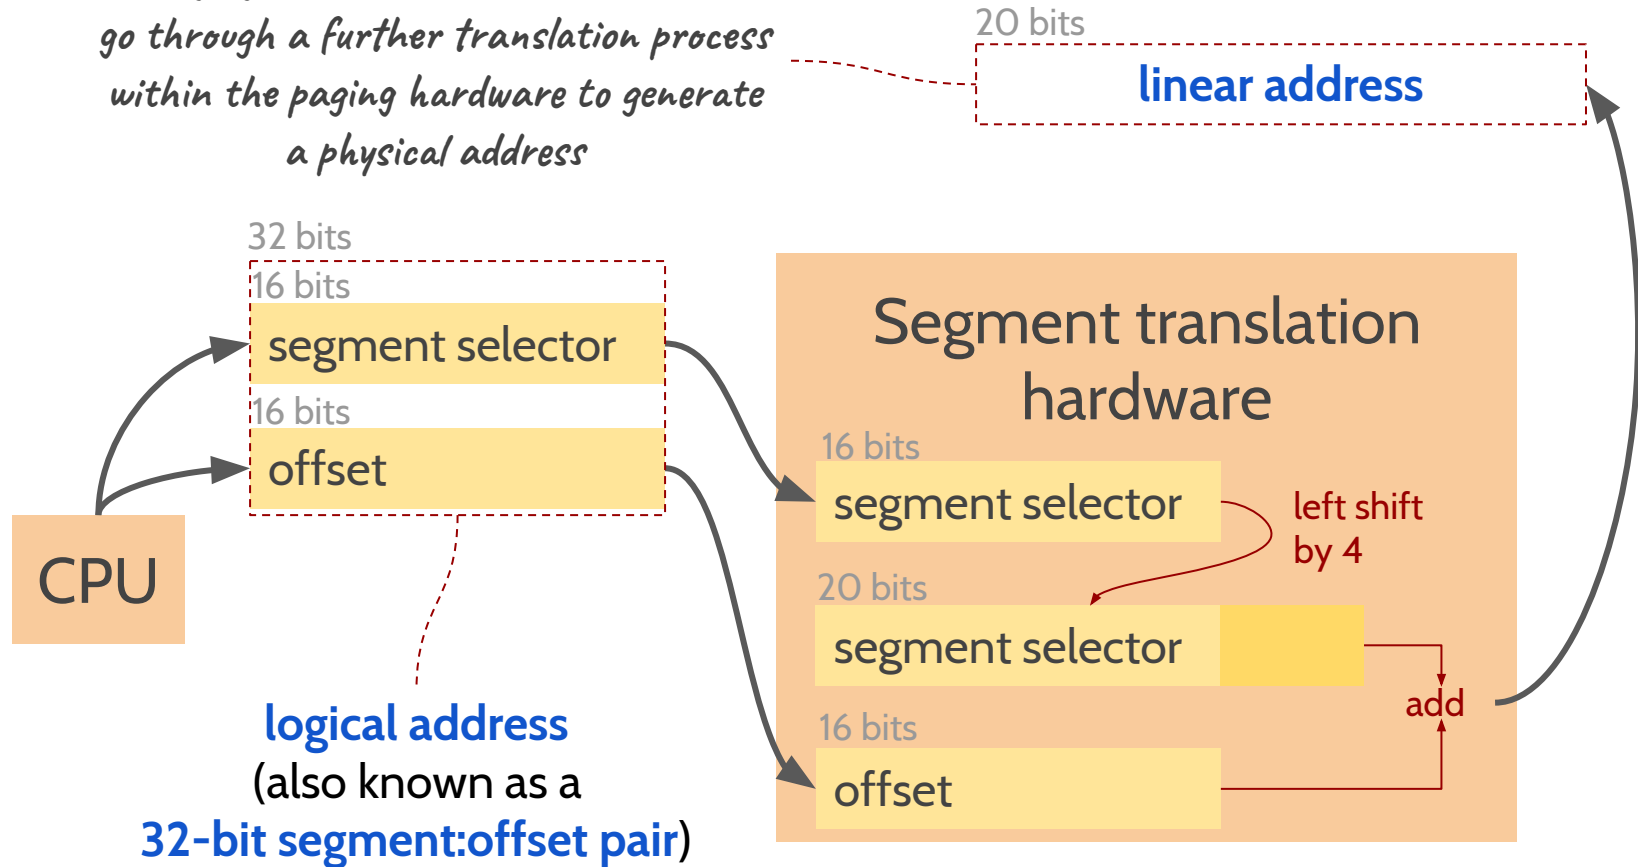
16 bits

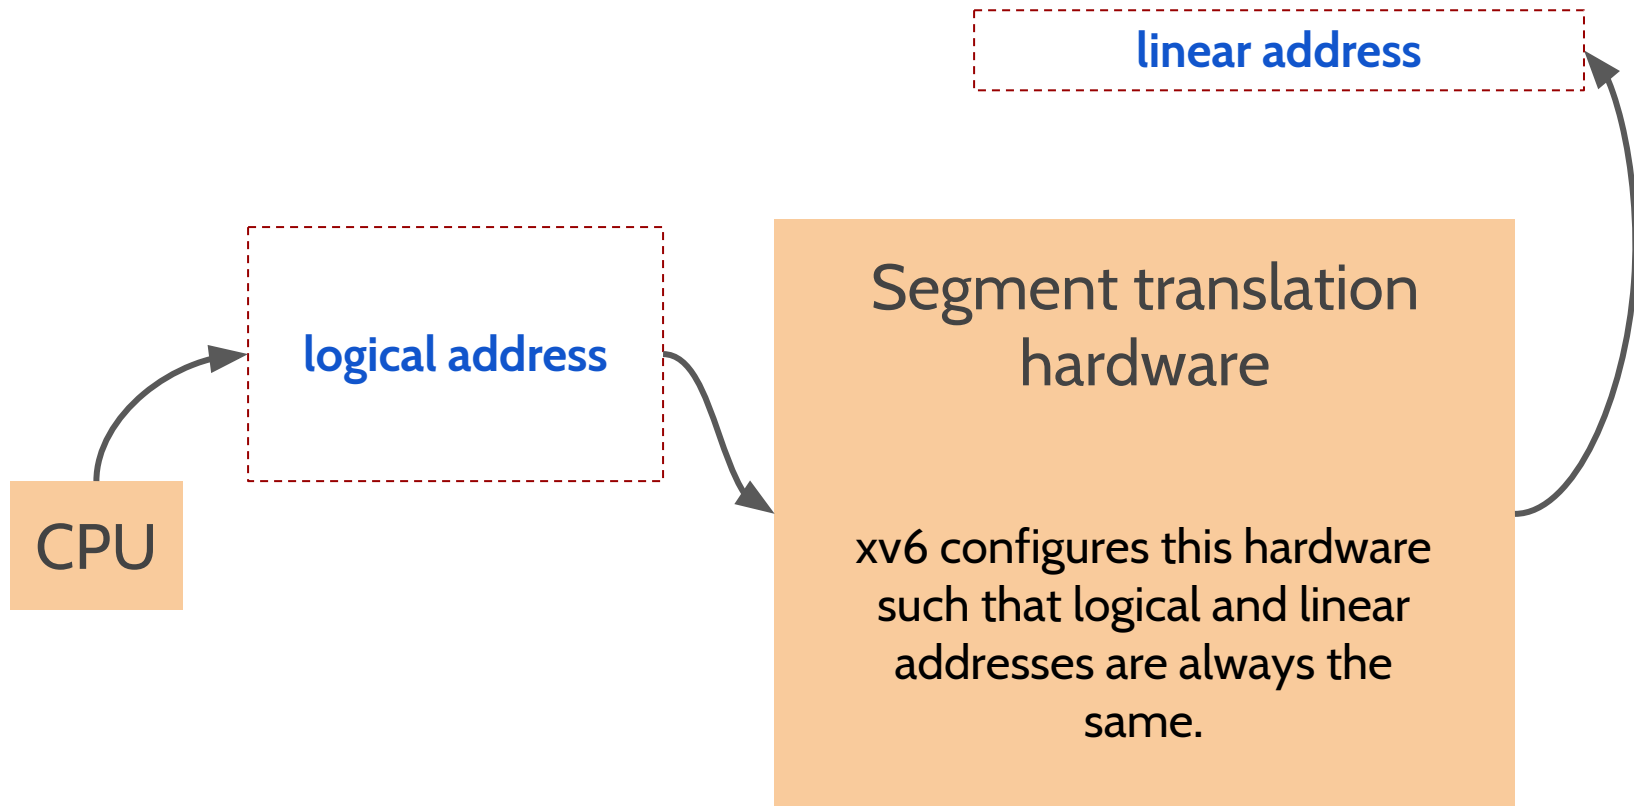
offset

add

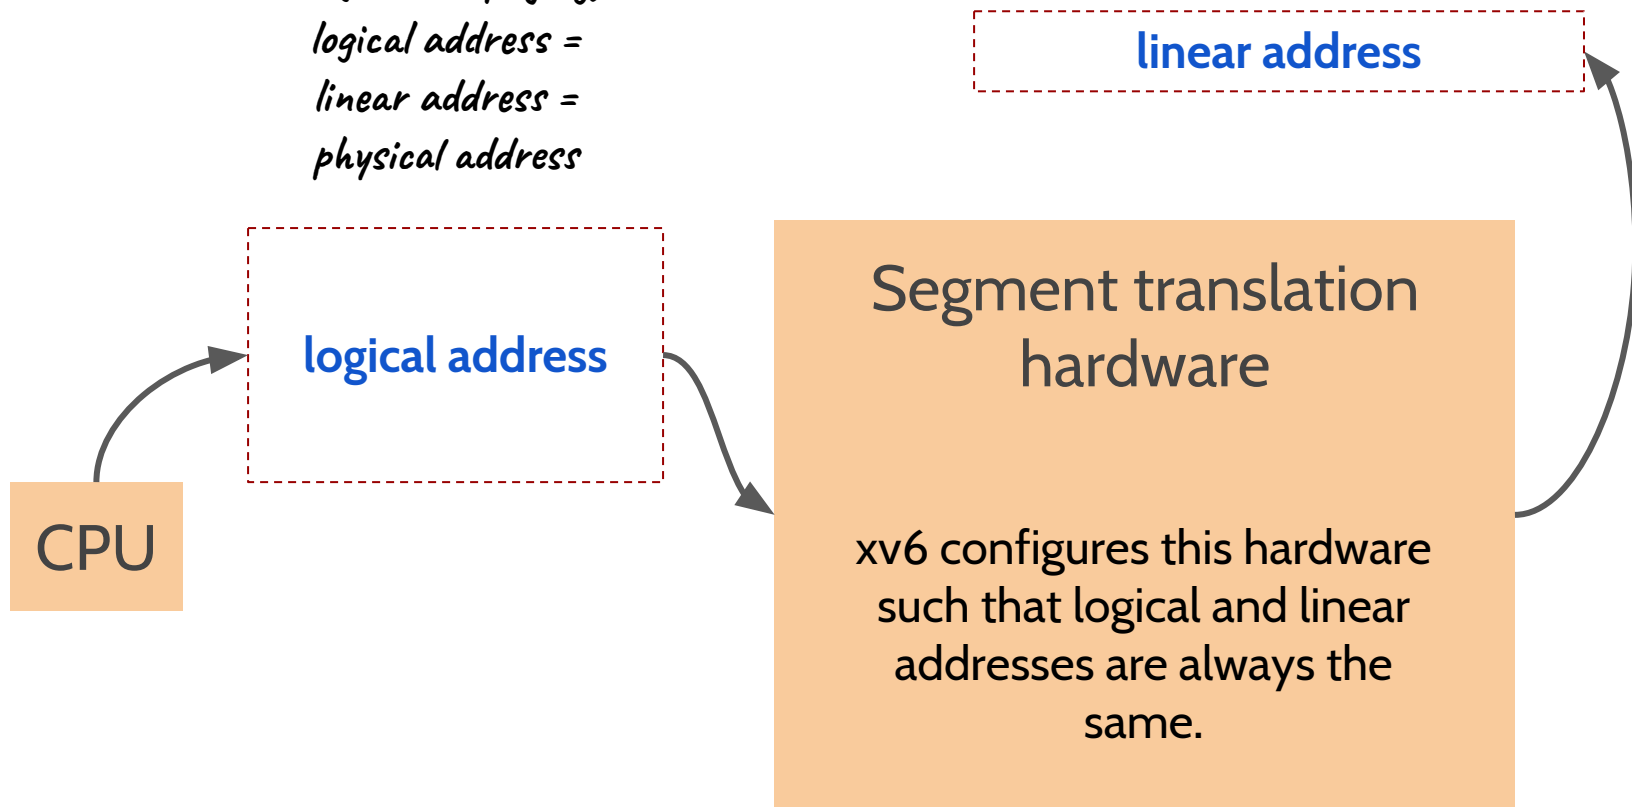


if paging is enabled, this address would go through a further translation process within the paging hardware to generate a physical address





*It follows (without paging) in xv6,
logical address =
linear address =
physical address*



Why the switch to Protected Mode?

with **20 bit addresses** the maximum size of addressable memory is 2^{20} bytes which is **only 1MB**.

In Protected Mode, we can address **32 bit** addresses, which allows us to address a memory of size **4GB**.

Paging hardware is also enabled in Protected Mode.

BIOS → xv6 Boot loader

executing bootasm.S

2. **Switch from real mode to protected mode.**
 - > To setup the **protected mode**, the boot loader sets up the **segment descriptor table**, with three entries. Each entry is [a base physical address, max virtual address limit (4GB), permission bits]

BIOS → xv6 Boot loader

executing bootasm.S

```
lgdt    gdtdesc
```

```
gdtdesc:
```

```
.word    (gdtdesc - gdt - 1)    # sizeof(gdt) - 1
```

```
.long    gdt                    # address gdt
```

setting up segment
descriptor table

the gdt table

the macros
used are
defined here

BIOS → xv6 Boot loader

executing bootasm.S

Assembly Tip:

All assembler directives begin with a period. ([Ref.](#))

```
lgdt    gdtdesc
```

```
gdtdesc:
```

```
.word    (gdtdesc - gdt - 1)    # sizeof(gdt) - 1
```

```
.long    gdt                    # address gdt
```

setting up segment
descriptor table

the gdt table

the macros
used are
defined [here](#)

BIOS → xv6 Boot loader

executing bootasm.S

2. **Switch from real mode to protected mode.**
 - > Now enable protected mode by setting the 1 bit in control register %cr0.
([Code](#)).

BIOS → xv6 Boot loader

executing bootasm.S

2. **Switch from real mode to protected mode.**
 - > To fully enable protected mode, we need to load a new value into %cs.
 - > Since %cs cannot be modified directly we setup %cs here: ([Code](#)), which results in %cs to refer to the code descriptor entry in gdt.

BIOS → xv6 Boot loader

executing bootasm.S

Assembly Tip:

Long Jump:

```
ljmp $0xfebc, $0x12345678
```

Use 0xfebc for the CS register and 0x12345678 for the EIP register. ([Ref.](#))

2. Switch from real mode to protected mode.

- > To fully enable protected mode, we need to load a new value into %cs.
- > Since %cs cannot be modified directly we `ljmp` ([code](#)), which results in %cs to refer to the code descriptor entry in gdt (a 32 bit code seg.)

BIOS → xv6 Boot loader

**We have completed transitioning from 16 bit (real) mode to
32 bit (protected) mode**

BIOS → xv6 Boot loader

To see details on how the 32 bit addresses are translated in the segmentation hardware in protected mode look [here](#).